

Distributing Parts of a Matrix from Master to Slaves Over Sockets with Multi-threading and CPU Affinity

Jadon C. Sacayanan

Abstract—This research examines the performance of a multi-threaded master-slave architecture for distributing parts of a matrix using sockets, both within a single PC and across multiple PCs. The study compares scenarios with and without CPU affinity. The focus is on the average runtime as the number of slaves increases. Results indicate that increasing the number of slaves reduces runtime significantly, with distributed computing across multiple machines showing the best performance.

Index Terms—parallel computing, multi-threading, CPU affinity, Z-score normalization, matrix computation

I. INTRODUCTION

Efficient data processing is crucial in high-performance computing, where optimizing computational tasks can significantly improve execution speed. When managing large datasets, data can be distributed to enable parallel processing, which may reduce processing time and improve overall performance.

A. Distributed computing

Distributed computing refers to programs that make calls to other address spaces, which may reside on the same machine or potentially on remote machines. This includes technologies such as remote object invocation, where tasks are distributed across different systems, allowing for the execution of processes on separate nodes and facilitating efficient resource utilization in a networked environment [1].

B. Socket programming

Socket programming is a way of connecting two nodes on a network to communicate with each other. One socket(node) listens on a particular port at an IP, while the other socket reaches out to the other to form a connection. The server forms the listener socket while the client reaches out to the server [2].

C. Thread

A thread is an abstract entity that represents the execution of a sequence of instructions within a program. It is a lightweight component consisting of essential elements like registers, a stack, and other associated data. A thread is the smallest unit of

execution, capable of being scheduled and executed independently on a CPU. Each thread can run independently, allowing multiple threads to execute in parallel, thereby improving the efficiency of the program [3].

D. CPU affinity

CPU affinity is a technique that binds a thread or process to a specific CPU core, preventing the operating system from migrating it to another core. CPU affinity is optimizing cache performance. When a thread is bound to a specific core, it can take advantage of the cache's contents, reducing cache misses and improving performance [4].

E. Cache awareness

When the processor needs to access data, it first checks its cache. If the data is found in the cache, it is called a cache hit, allowing for fast retrieval. However, if the data is not present, a cache miss occurs, requiring the processor to fetch the data from main memory, which is significantly slower.

The processor does not work by retrieving a single byte of data, but instead fetches a larger block of memory called a cache line. By loading an entire cache line at once, the processor increases the likelihood that future memory accesses will be served from the cache rather than requiring additional slow memory fetches. This reduces the frequency of cache misses and enhances overall processing efficiency [5].

F. Temporal locality

Programs has the tendency to repeatedly use the same data items during their execution. This principle is the basis for caching and suggests that it is beneficial to store frequently accessed instructions or data items nearby for future use [6].

II. OBJECTIVES

This research activity aims to analyze the performance and efficiency of distributing matrix computations across multiple processes using socket communication. Specifically, it aims to:

- 1) Implement a master-slave architecture that distributes submatrices to multiple processes via socket communication.
- 2) Investigate the impact of varying the number of slave processes (t) on the execution time of matrix distribution tasks.

Presented to the Faculty of the Institute of Computer Science, University of the Philippines Los Baños in partial fulfillment of the requirements for the Degree of Bachelor of Science in Computer Science

- 3) Compare the performance when all processes are run on a single machine versus when processes are distributed across multiple physical machines.
- 4) Evaluate the effect of core-affinity on communication efficiency when processes execute on the same physical machine.

III. MATERIALS AND METHODS

A program was developed to distribute parts of an $n \times n$ matrix from a master to multiple slave machines using sockets. For each matrix size n , the distribution process was executed three times to calculate the average runtime. To ensure accurate measurements, all tests were conducted under controlled conditions with no other applications running, minimizing external influences on network and system performance.

A. The program

The program is written in C and begins by prompting the user to input three values: n , the size of the square matrix; p , the port number used for socket communication; and s , a value that determines the role of the program, where 0 indicates a master and 1 indicates a slave. These inputs define how the program will behave during the matrix distribution process.

B. Master

When the program is run as a master, it generates an $n \times n$ random square matrix. The matrix is then divided into submatrices, each with n columns and approximately n/t rows, where t is the number of slaves. If n is not exactly divisible by t , the remaining rows are distributed among the first few slaves, with each receiving at most one additional row. Rather than creating separate matrices, the master identifies the start and end row indices for each submatrix, referencing the original matrix directly.

Once all submatrices are determined, the master establishes socket connections to the slave machines. It then sends each submatrix, number by number, to its assigned slave. After transmission, the master waits for an acknowledgment from each slave to confirm that the submatrix was received successfully.

C. Slave

When the program is run as a slave, it waits for a connection from the master. Upon connection, the slave receives the submatrix data, transmitted number by number, and stores it in a dynamically allocated array sized according to the expected number of rows and columns. After successfully receiving the entire submatrix, the slave sends an acknowledgment back to the master to confirm receipt. Once the communication is complete, the socket connection is closed.

D. Threading

The program uses the pthread library to create threads. In both the master and slave roles, each thread is responsible for handling a single socket connection, starting from establishing

the connection (sending or listening) until the communication is closed. This design allows concurrent handling of multiple connections, enabling parallel distribution and reception of submatrices.

E. CPU affinity

The program determines the number of available cores using the sysconf function and assigns each thread to a specific core using the CPU affinity feature. To ensure system stability and prevent resource contention, one core is reserved for main system processes, such as the operating system scheduler, background services, and essential system tasks. Each thread is pinned to a dedicated CPU core using the CPU set function, preventing the operating system from migrating threads between cores and reducing interference from other system functions. This avoids the overhead of thread migration and maximizes program performance.

F. Time measurement

The program uses the sys/time.h header to measure execution time by capturing a timestamp with gettimeofday() immediately before establishing a connection. A second timestamp is recorded after the socket is closed. The elapsed time is then calculated by determining the difference between these two timestamps. This time measurement only includes the communication phase and does not account for the preparation of the threads, such as setting up their arguments.

G. Distributed computing

The program utilizes four machines, running up to four instances of the slave program on each machine. Different configurations were used depending on the value of n . For $n=2$, two machines were used, each running one instance of the slave. For $n=4$, four machines were used with one instance of the slave on each. For $n=8$, the same four machines were used, each running two instances of the slave. Finally, for $n=16$, all four machines ran four instances of the slave each.

It also utilized a personlaized one-to-many broadcast mechanism, where the master node efficiently sends data to multiple slave nodes simultaneously. This approach minimizes communication overhead by leveraging non-blocking socket operations, ensuring that the master can continue transmitting data to other nodes without waiting for acknowledgments from previously contacted nodes.

H. System information

The program was executed on an Acer desktop equipped with an Intel® Core™ i7-8700 processor (3.20 GHz) and 16GB of memory, running Ubuntu 22.04.5 LTS. Both the master and all slave machines used the same hardware and software configuration.

IV. RESULTS AND DISCUSSION

The program was executed independently for three iterations for each setup, recording the runtime or elapsed time for each execution.

TABLE I

RUNTIME EXECUTIONS FOR DIFFERENT N AND T VALUES IN SECONDS
FOR DISTRIBUTED COMPUTING ON DIFFERENT MACHINES

n	t	Time Elapsed (seconds)			Average Runtime
		Run 1	Run 2	Run 3	
20000	2	158.073970	161.382508	154.220397	157.892292
20000	4	92.240668	88.906241	94.108334	91.751748
20000	8	66.350505	68.771319	64.112398	66.411407
20000	16	49.856434	50.602118	47.981374	49.479975
25000	2	258.821180	267.093215	252.543018	259.485804
25000	4	143.290254	140.113850	146.001709	143.135271
25000	8	110.955387	108.144267	113.589460	110.896371
25000	16	72.869340	74.302138	70.004250	72.391909
30000	2	371.310715	377.498020	365.918605	371.575780
30000	4	179.321567	183.927205	175.203491	179.484088
30000	8	161.082212	164.189404	158.440373	161.237330
30000	16	138.598874	136.221670	141.785209	138.868584



Fig. 1. Runtime Executions for n=20000 in seconds for Distributed Computing on Different Slaves using Different Machines

A. Distributed computing with slaves running on different machines

Table I shows the average runtime for distributing an $n \times n$ matrix using t threads across different slaves running on different machines. The results indicate that as the number of slaves increases, the runtime decreases significantly.

For a square matrix of size 20,000, the slowest runtime of approximately 157.89 seconds is observed when using 2 slaves. In contrast, the fastest runtime of approximately 49.48 seconds is achieved when using 16 slaves. This trend is visualized in Figure 1, where the runtime consistently decreases as the number of slaves increases.

B. Distributed computing with slaves running on a single machine without core affinity

Table II shows the average runtime for distributing an $n \times n$ matrix using t threads across different slaves running on a single machine without core affinity. The results indicate that as the number of threads increases, the runtime initially decreases but then starts to increase again after $t=8$.

For a square matrix of size 20,000, the fastest runtime of approximately 106.85 seconds is observed when using 4 threads. However, when the number of threads is increased to 16, the runtime increases significantly to approximately 207.86 seconds. This trend is visualized in Figure 2, where the runtime

TABLE II

RUNTIME EXECUTIONS FOR DIFFERENT N AND T VALUES IN SECONDS
FOR DISTRIBUTED COMPUTING ON DIFFERENT SLAVES USING A SINGLE
MACHINE WITHOUT CORE AFFINITY

n	t	Time Elapsed (seconds)			Average Runtime
		Run 1	Run 2	Run 3	
20000	2	164.150009	166.822194	160.485729	163.819311
20000	4	106.496193	109.841512	104.219304	106.852336
20000	8	111.433945	108.756837	113.275409	111.155397
20000	16	209.123130	202.578341	211.892084	207.864518
25000	2	257.352387	263.091982	254.305673	258.250014
25000	4	174.123655	176.452399	171.009812	173.861955
25000	8	169.091237	173.208420	165.877304	169.392320
25000	16	344.907369	337.122408	352.610193	344.879990
30000	2	369.370952	362.041733	375.008091	368.806925
30000	4	254.110043	260.344552	247.812063	254.088886
30000	8	247.671020	251.493800	241.520991	246.895270
30000	16	512.400644	523.108477	505.287190	513.598770

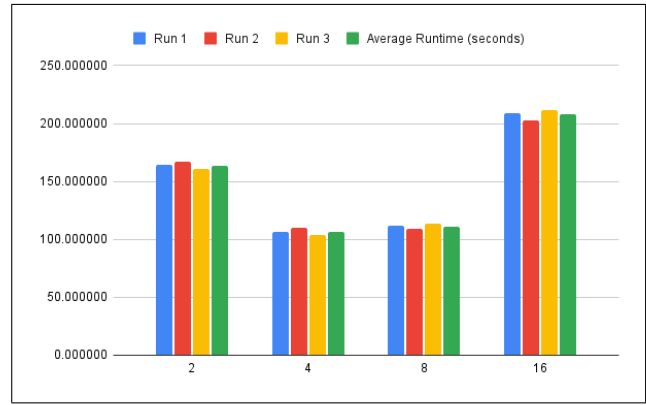


Fig. 2. Runtime Executions for n=20000 in seconds for Distributed Computing on Different Slaves using a Single Machine without Core Affinity

decreases initially but rises sharply as the number of threads exceeds the number of physical cores.

C. Distributed computing with slaves running on a single machine with core affinity

Table III shows the average runtime for distributing an $n \times n$ matrix using t threads across different slaves running on a single machine with core affinity. The results indicate that as the number of threads increases, the runtime decreases significantly, even for higher thread counts.

For a square matrix of size 20,000, the fastest runtime of approximately 46.55 seconds is observed when using 16 threads. This trend is visualized in Figure 3, where the runtime consistently decreases as the number of threads increases.

V. CONCLUSION AND FUTURE WORK

Parallel computing enhances program performance by distributing workloads across multiple nodes, enabling efficient utilization of computational resources. The results validate that increasing the number of nodes generally reduces runtime, as observed in both distributed and single-machine setups.

For distributed computing across multiple machines, increasing the number of slaves significantly improved performance, with runtimes decreasing as the number of slaves

TABLE III

RUNTIME EXECUTIONS FOR DIFFERENT N AND T VALUES IN SECONDS
FOR DISTRIBUTED COMPUTING ON DIFFERENT SLAVES USING A SINGLE
MACHINE WITH CORE AFFINITY

n	t	Time Elapsed (seconds)			Average Runtime
		Run 1	Run 2	Run 3	
20000	2	165.097044	168.201223	162.89451	165.3975923
20000	4	160.935094	159.723	161.889234	160.8491093
20000	8	101.691484	104.229843	99.875621	101.932316
20000	16	46.703392	47.115298	45.834002	46.55089733
25000	2	268.059515	270.64312	265.202891	267.9685087
25000	4	255.063759	254.829371	258.140574	256.0112347
25000	8	159.475293	162.008734	158.027191	159.8370727
25000	16	75.655035	76.308295	74.981402	75.648244
30000	2	392.014359	390.202773	395.451002	392.5560447
30000	4	342.224166	345.11498	341.338267	342.892471
30000	8	232.154717	229.119456	234.203887	231.82602
30000	16	98.284437	99.087141	97.652309	98.34129567

- [5] Tennessee Tech University. (2025, Feb.) Cache awareness. [Online]. Available: <https://tinyurl.com/2aj2x8tz>
- [6] B. Jacob, D. Wang, and S. Ng, *Memory systems: cache, DRAM, disk*. Morgan Kaufmann, 2010.




Fig. 3. Runtime Executions for n=20000 in seconds for Distributed Computing on Different Slaves using a Single Machine with Core Affinity

increased. This highlights the benefits of leveraging distributed systems for large-scale computations.

On a single machine, the use of core affinity showed mixed results. While core affinity aims to optimize cache utilization and reduce thread migration overhead, the implementation without core affinity was faster in some cases. This suggests that the system's scheduling algorithm may already optimize thread placement effectively, and the manual implementation of core affinity might introduce additional overhead. When the number of threads exceeds the number of available cores, the implementation without core affinity experiences a significant increase in runtime.

In conclusion, optimizing performance in parallel computing requires an approach that considers hardware capabilities, memory access patterns, and system-level optimizations. This study provides a foundation for further exploration of efficient parallel and distributed computing techniques.

REFERENCES

- [1] J. Waldo, G. Wyant, A. Wollrath, and S. Kendall, "A note on distributed computing," in *International Workshop on Mobile Object Systems*. Springer, 1996, pp. 49–64.
- [2] GeeksforGeeks. (2025, Apr.) Socket programming in c. [Online]. Available: <https://www.geeksforgeeks.org/socket-programming-cc/>
- [3] B. Lewis and D. J. Berg, "Pthreads primer," *Sun Microsystems Inc*, 1996.
- [4] R. Love, "Kernel korner: Cpu affinity," *Linux Journal*, vol. 2003, no. 111, p. 8, 2003.